



Expert Developer Recommendation Using Very Large Datasets

Project Presentation

Saim Mehmood
Tilemachos Pechlivanoglou



Project goals

We want to create a tool that:

- can be used to search for expert developers
- provides search criteria for developers' skills and experience
- bases results on actual contributions
- uses a very large dataset of possible experts (20+ million)
- search queries are reasonably fast
- can be extended to include more criteria
 - *(code quality, bugs introduced, etc.)*



Example search

- Looking for a developer who:
 - *Has 5 years expertise in C++*
 - *Has contributed to an Apache project*
- Or:
 - *Worked recently with Python*
 - *Is familiar with (has used) the Flask framework*
 - *Commits often*

Pre-Processing



BigQuery and big troubles



The dataset available

- The dataset:
 - *20+ million developers*
 - *215 million commits*
 - *2.3 billion files*
 - *3+ TB of data*
- Sample dataset available
 - *used during development*
 - *500 times smaller (0.2%)*



Some dataset issues

- Some columns available in sample missing in full
 - *Workarounds needed*
- Available online only
 - *Can't directly download data above a few MB*
 - *Too big to process locally anyway*
 - *Preprocessing needed to reduce size*



BigQuery

Google's BigQuery is perfect for preprocessing

- Allows running SQL queries on dataset
 - *Traditional and extended SQL supported*
- Distributed processing, optimizations for big data
 - *Mostly transparent to the user*



Data reduction (1)

Step 1: Throw away useless data

- We focus on developer commit data only
- Other information is discarded:
 - *file contents may be used in the future*
 - *hashes etc. not useful to the project*
- Data reduction: 3TB → 100GB !!





Data reduction (2)

Step 2: Associate file extensions with programming languages/frameworks

- Find files altered in every commit
- Match file extensions to languages
 - *Using external list*
 - *Throw away file path data*
- Data reduction: 100GB → 51GB


Extension		Language/ Framework
.cpp	→	C++
.java	→	Java
...	→	...



Data reduction (3)

Step 3: Condense time data

- For every developer skill:
 - *Aggregate commits*
 - *Store first, last commit*
 - *Store nr. of commits, files*
- Data reduction: 51GB → 1.1GB !!






Data reduction (4)

Step 4: Small optimizations

- Reducing forked repo names
 - *e.g. torvalds/linux* → *linux*
- Representing as JSON, not CSV
 - *nested data, less duplication*
- Data reduction: 1.1GB → 600MB



Elasticsearch

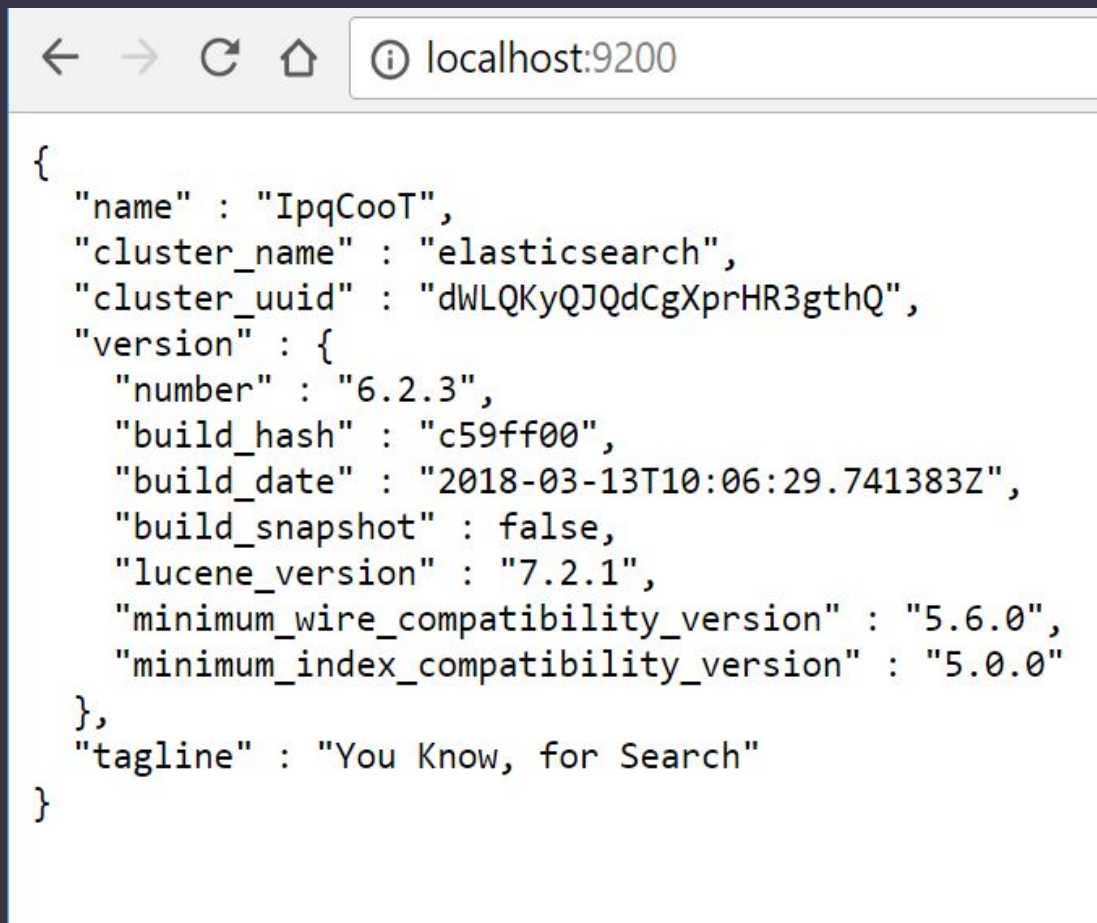
Performing Queries



About the tool

- Elasticsearch is a **highly scalable** open-source full-text search and analytics engine
- Allows you to store, search, and analyze big volumes of data quickly and in near real time
- Used as the underlying engine/technology that powers applications with complex search features

14

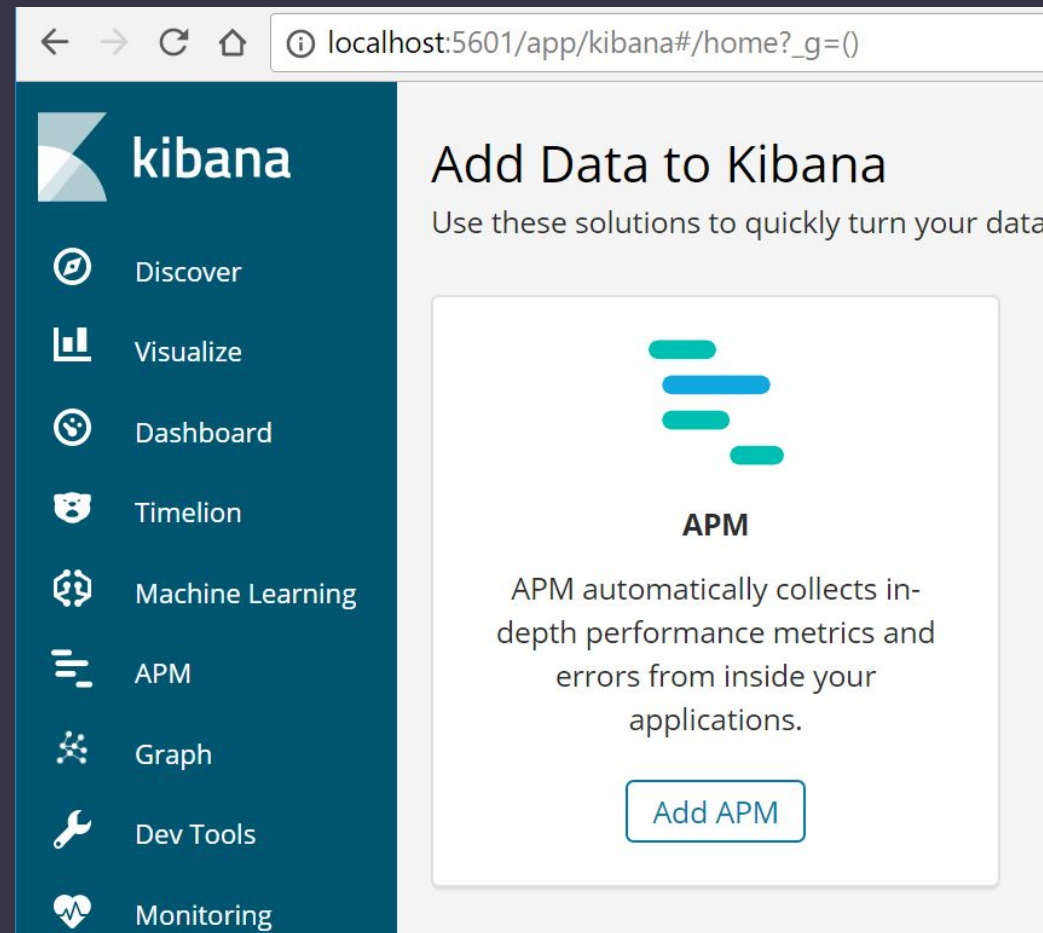


```
{
  "name" : "IpqCooT",
  "cluster_name" : "elasticsearch",
  "cluster_uuid" : "dWLQKyQJQdCgXprHR3gthQ",
  "version" : {
    "number" : "6.2.3",
    "build_hash" : "c59ff00",
    "build_date" : "2018-03-13T10:06:29.741383Z",
    "build_snapshot" : false,
    "lucene_version" : "7.2.1",
    "minimum_wire_compatibility_version" : "5.6.0",
    "minimum_index_compatibility_version" : "5.0.0"
  },
  "tagline" : "You Know, for Search"
}
```

Elasticsearch Server

Initial Working

>.\bin\elasticsearch



Kibana

>.\bin\kibana



Steps

- Elasticsearch requires indexing for every object (i.e. developer) inside the JSON file
- We used text editor for indexing (avoiding the overhead of Kibana)
- Uploading data into Elasticsearch
 - `curl -H "Content-Type: application/json" --user elastic:elastic`
`-XPOST "localhost:9200/developer/email/_bulk?pretty"`
`--data-binary @developers.json`



Sample Query

- Simple query to fetch developers by name:

- `curl -H "Content-Type: application/json" --user elastic:elastic`
`-XGET "localhost:9200/_search?pretty"`
`-d '{"query":{"match":{"developer": "Aaron"}}}'`



Nested Query

- Nested query to refine search criteria:
 - `curl -H "Content-Type: application/json" --user elastic:elastic`
`-XGET "http://localhost:9200/_search?pretty"`
`-d '{"query": { "nested": {"path": "languages",`
`"query": { "bool" : { "must": [{ "match": { "languages.name" : "Objective-C"} },`
`{ "range": { "languages.duration" : {"gt": 3567688 } } }] } } } } } }`




Live Demo



Show time!



The whole system



Conclusion

Things we learned



Future Goals

- Ranking of developers based on Elasticsearch scoring mechanism
- Extending tool to provide search of projects and repositories
- Processing code commits to find Code Quality Metrics specific to a project or developer
- Evaluating tool on diverse data sets



Potential future steps

With this project established, the recommendation algorithm could be extended

- Detailed analysis of code
 - *developer's code wasn't immediately replaced*
 - *code didn't introduce bugs (bugfix comments)*
- More complex requirements
 - *Search for group of experts who cooperated in the past*



Issues

- Financial Constraint: BigQuery 1TB of processing/month for free
- Full data set doesn't have files changed in commit
 - *Impossible to exactly determine user skills*
 - *Approximation based on repo languages*
 - *Can be mitigated by analyzing file contents*
- Most issues can be resolved with time and/or money



Project goals

We want to create a tool that:

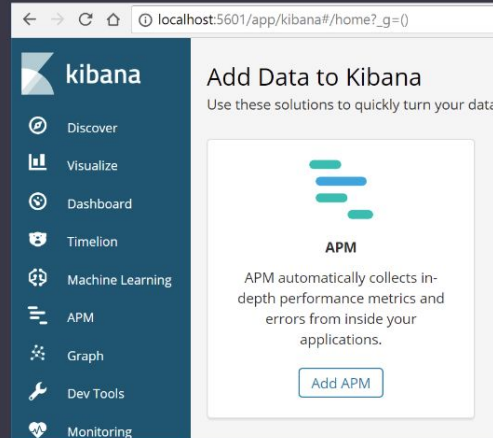
- can be used to search for expert developers
- provides search criteria for developers' skills and experience
- bases results on actual contributions
- uses a very large dataset of possible experts (20+ million)
- search queries are reasonably fast
- can be extended to include more criteria
 - (code quality, bugs introduced, etc.)

```

{
  "name" : "IpqCooT",
  "cluster_name" : "elasticsearch",
  "cluster_uuid" : "dWLQKyQJQdCgXprHR3gthQ",
  "version" : {
    "number" : "6.2.3",
    "build_hash" : "c59ff00",
    "build_date" : "2018-03-13T10:06:29.741383Z",
    "build_snapshot" : false,
    "lucene_version" : "7.2.1",
    "minimum_wire_compatibility_version" : "5.6.0",
    "minimum_index_compatibility_version" : "5.0.0"
  },
  "tagline" : "You Know, for Search"
}


```

Elasticsearch Server



Kibana

The whole system



Nested Query

- Nested query to refine search criteria:

```

- curl -H "Content-Type: application/json" --user elastic:elastic
-XGET "http://localhost:9200/_search?pretty"
-d '{"query": { "nested": { "path": "languages";
"query": { "bool": { "must": [{ "match": { "languages.name" : "Objective-C" } },
{ "range": { "languages.duration" : { "gt": 3567688 } } } ] } } } } }'

```

Thank you! Questions?

